

The ASGS Operators Guide

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
1.0	June 2012, Seahorse Coastal Consulting		JF

Contents

1	Executive Summary	1
2	System Structure	1
2.1	Reference Information	3
2.2	Dynamic Input Data	3
2.3	Input File Production	3
2.4	ADCIRC Asymmetric Vortex Model Used During Irene	5
2.4.1	Dynamic Asymmetric Holland Model	5
2.5	Output Visualization, Publication, and Interaction	6
3	Configuration	7
3.1	Fundamental	7
3.2	Source File Paths	8
3.3	Physical Forcing	9
3.4	Computational Resources	10
3.4.1	Wall Clock Time Estimates	11
3.4.2	Dynamic Configurability	11
3.5	External Data Sources	12
3.5.1	Tropical Cyclones	12
	Testing	13
	Interactions	13
3.5.2	Background Meteorology	13
3.5.3	River Flux	14
3.6	Input Files and Templates	14
3.7	Storm Ensemble	16
3.8	Output Files	17
3.9	Notification	18
3.10	Post Processing and Publication	20
3.11	Archiving	21
4	Execution	21
5	Testing	21
5.1	Setting Up the Mock NHC Files for Testing	21
5.2	Configuring the ASGS for Test Advisories Stored on Local Filesystem	22
5.3	Actually Running ASGS from Advisories Stored in Local Files	22

6	Performance During Hurricane Irene	23
6.1	North Carolina Physical Data	23
6.2	Dynamic Input Data	23
6.3	Internally Generated Input Files	24
6.4	Output Visualization and Interaction	24
7	Performance During the Deepwater Horizon Event	24
8	Refs	24
A	This Document	24

1 Executive Summary

The ADCIRC Surge Guidance System (ASGS) is a software system for generating storm surge and wave guidance from ADCIRC + SWAN in real time on high resolution grids. The Coastal Emergency Risks Assessment (CERA) web visualization application provides an interactive web medium for examination of results by official clients. During a tropical cyclone event, the ASGS constructs meteorological forcing from a parametric wind / pressure model using storm parameters extracted from National Hurricane Center (NHC) Forecast Advisories. During nor'easters or for day-to-day generation of tidal forecasts and other results, the system uses gridded wind and pressure fields (e.g., NCEP's NAM model) as input forcing. In both cases, hydrologically-driven river forecast data from NSSL are used to account for the effects of precipitation and upland river flooding.

2 System Structure

An overview of the modular structure of the ASGS is shown in Figure 1. The figure indicates the separation of System components into the following categories:

1. reference information (in purple), which specify the system configuration and physical parameter data used in the simulation;
 2. dynamic input data (in red) that varies in time and must be downloaded from external data sources for every forecast cycle;
 3. input file production (in green), which implements the system behavior specified by the Operator in the system configuration files; and
 4. output visualization, publication, and interaction (in blue) for use by clients and end users.
-

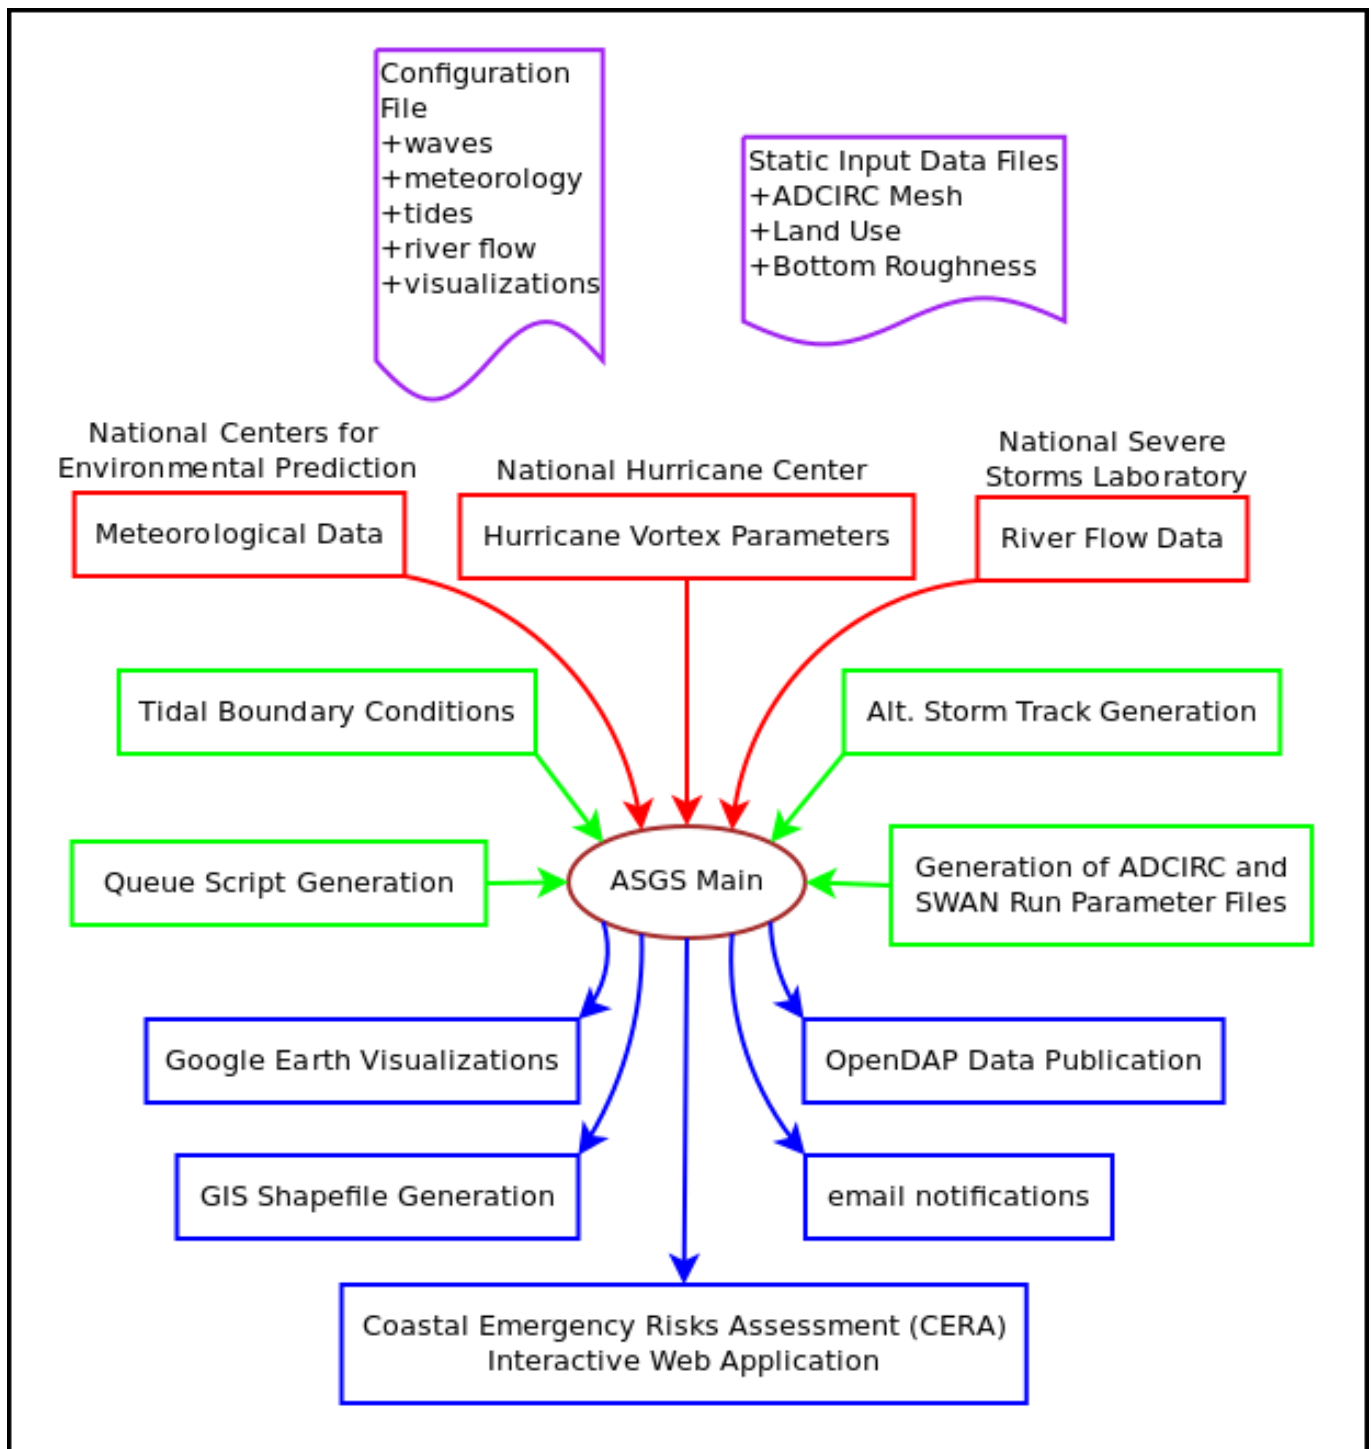


Figure 1: The overall structure of the ASGS divides the various features into their own modules. The configuration and static physical data that are common to all simulation runs are shown at the top in purple, the dynamic data acquisition modules are shown in red, the internal data processing and input file generation are shown in green, and output modules are shown in blue at the bottom. Arrows are conceptual and indicate data flow.

These categories and associated modules will be described in greater detail below.

2.1 Reference Information

The reference information includes a dynamic specification of system behavior in the configuration file as well as static physical data, embodied in the simulation input files ADCIRC. The system configuration consists of a single file, and all of the features and behavior of the various components can be controlled from this single configuration file. This file is used to activate or deactivate the various types of physical forcing, including tropical cyclone meteorology, ordinary meteorology, wave simulation and coupling, tides, and river flow input. It is also used to specify a wide variety of other settings, including (for example) the type of computer that the system is running on, the number of processors to use in parallel execution, the name of the ADCIRC mesh file for the domain of interest, the number of storms in an ensemble and their characteristics, the types of output products to generate, the email addresses of officials that should be notified when results are ready, and many others.

The simulation input files represent a purely static set of physical data that are used in the simulations. These data include the ADCIRC mesh (domain discretization), the bathymetry and topography of the domain, the spatially varying Manning's n value, and the directional wind roughness lengths and canopy effects derived from land cover data. The static data also include the tidal constituents to be included (if any), convergence controls and solution parameters for SWAN and ADCIRC, the names and locations of point recording stations for location-specific output data, and other data related to the internal procedures of the simulation codes.

2.2 Dynamic Input Data

The system is capable of downloading and preparing dynamic meteorological data from two sources: the National Hurricane Center's (NHC) Forecast/Advisories, and the National Centers for Environmental Prediction's (NCEP) North American Mesoscale (NAM) model, depending on the data source selected by the Operator. Furthermore, the system has a module for downloading river boundary flow data from the National Severe Storms Laboratory and preparing it for use in ADCIRC.

For tropical cyclone events, the NHC Forecast/Advisories are downloaded by the system as soon as they are issued, and the relevant storm parameters (including storm positions, maximum wind speeds, and isotach radii) are parsed into a format that ADCIRC can use to generate wind fields using an internal asymmetric vortex model. For ordinary meteorology, including nor'easters and other systems, the system downloads regularly gridded meteorological fields from the NAM model from the NCEP ftp site. The gridded meteorological fields are extracted from the grib2 format that NCEP uses, reformatted and written to OWI formatted files that ADCIRC reads.

Once meteorological data have been obtained, the system downloads the latest river flow data from the National Severe Storms Laboratory (NSSL) ftp site. These data are in a file format that ADCIRC can read natively, and are set up to correspond to a particular ADCIRC mesh. The module that downloads these data must splice a series of these files together to match the date and time range implied by the meteorological data that have already been obtained.

2.3 Input File Production

Once the Operator has created a configuration file and assembled the static physical input data for a particular instance of the system and initiates startup, the system reads its configuration file and tries to determine its state: whether it is starting "hot" with a simulation that is already in progress, or "cold", where a new ADCIRC or ADCIRC+SWAN simulation must be started from scratch (see Figure 2). It also makes a determination about the state of its input files; specifically, whether they have already been decomposed for parallel execution. If the system must start from a "cold" state, it creates input files and executes a hindcast simulation to warm up the simulation and prepare it for the cyclical nowcast/forecast production phase (described below). Accordingly, this hindcast run is set so that it writes a hotstart file at the very end to represent the state of the warmed-up simulation. In addition, if the input files have not been prepared for parallel execution, the system decomposes them for the specified number of processors. If the hindcast and/or decomposition tasks are required, they must only be performed once at the start of the system's execution (processes that occur at most one time during an execution of the system are shown in blue in figure 2).

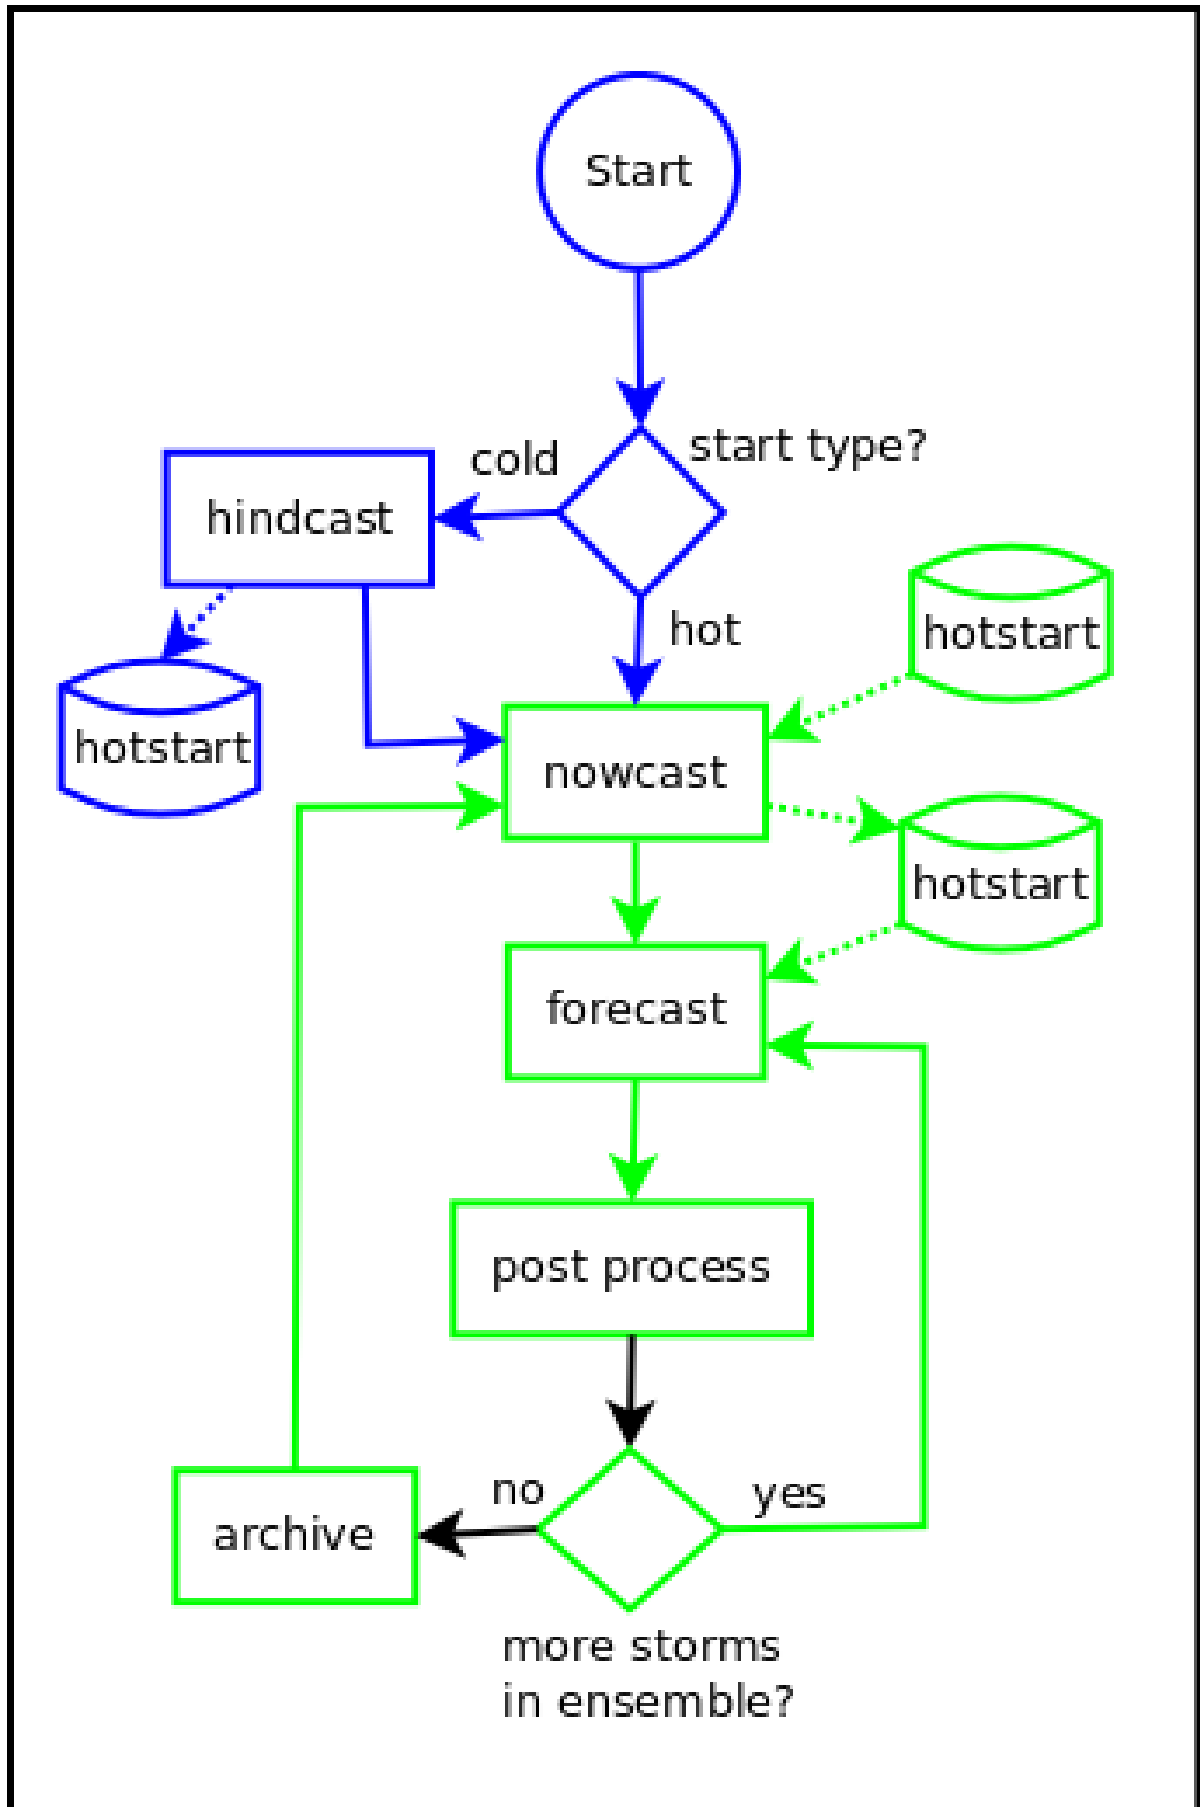


Figure 2: The overview of the logic of the ASGS divides the one-time start up process for an initially "cold" simulation state in blue, with the bulk of the time being spent in the nowcast/forecast cycle, indicated by green.

The nowcast occurs next if the simulation state is already "hot" when the system starts up, or if the system has already performed the decomposition and warm up phases. The first step in the nowcast is to determine if there are new meteorological data available, by contacting the associated external website or ftp site, checking the timestamps that are available there, and comparing them with the current simulation time. If there are data available that are more recent than the current simulation time, a new cycle is deemed to have begun, and those data are downloaded. The system downloads the data it requires to cover the time period between its current simulation time and the most recent data that are available files that are available. After the meteorological data have been acquired, river data are acquired in the same way, if they have been specified.

Once the external data have been acquired, the input files for the ADCIRC simulation code (and the SWAN simulation code if wave forcing has been specified by the Operator in the system configuration file) are constructed using the time range of the meteorological forcing files and the various configuration parameters in the system configuration file. This includes tidal boundary conditions, output file formats and frequency of production of output files, and locations for point recording of output for comparison with tide gages or meteorological data collection equipment. The nowcast control file(s) for ADCIRC and SWAN are set to write a hotstart file at the end of the simulation for use in starting the forecast, as well as a future nowcast. The last file written during the nowcast phase is the queue script that will be used to submit the job to the high performance computer's queueing system.

Once the nowcast is complete, the system acquires the data required for one or more forecasts. These data will already be present in the case of a tropical cyclone, since the NHC forecast/advisory has already been parsed. For NAM forcing, the NAM forecast data are downloaded from NCEP and converted to OWI using the same technique as described for the nowcast. If river flux forcing was specified, the river flux data are downloaded and formatted in a similar manner to the nowcast. The control files are constructed to cover the time period implied by the forecast meteorological data, but are not set to write a hotstart file at the end. When the forecast is complete, the system executes post processing (described in the next section) and archives the data as required.

These forecast process described above is applied to each member of the forecast ensemble until all ensemble members have been completed, at which point the system goes back to looking for new meteorological data for its next nowcast.

2.4 ADCIRC Asymmetric Vortex Model Used During Irene

The application of meteorological forcing presented a challenge for operational storm surge forecasts because of the need for timely availability of high resolution input data. The most accurate data-assimilated meteorological fields from the H*Wind project were only available for nowcast and hindcast times, and even then were not available until several hours had passed after a corresponding hurricane advisory had been issued from the NHC. In addition, regularly gridded meteorological data from models such as the North American Mesoscale (NAM) model had a relatively coarse grid resolution in comparison to the unstructured ADCIRC mesh (12km grid resolution vs 30m minimum mesh resolution).

In contrast, parametric wind models produce comparable storm surge in many cases (Houston et al, 1999; Mattocks et al 2006). They also had the following advantages: (1) they require a comparatively tiny quantity of input data; (2) they could be coded as fast subroutines that run in-process; (3) they could provide wind stress and barometric pressure at arbitrary locations.

As a result of the advantages of parametric wind models, the Holland model (Holland, 1980) was selected as the basis of the wind speed and pressure field. Modifications and additions were made to the published model to account for the dynamic changes in the hurricane parameters along the hurricane's track, as well as its adaptation to the data available from the NHC advisory, as described below. This modified model was referred to as the Dynamic Asymmetric Holland Model.

2.4.1 Dynamic Asymmetric Holland Model

The hurricane advisory contained at least the following information for each forecast increment: date, time, latitude and longitude of the center of the storm, and maximum observed wind speed at 10m with a 1 minute sampling interval. In addition, the distance to isotachs at various wind speeds were sometimes provided in one or more storm quadrants.

The vortex model was set up to operate in a two step process. The first step was the use of isotach data from the forecast/advisory to calculate a representative radius to maximum winds (Rmax) in each of the four storm quadrants, for each forecast increment, before the actual ADCIRC run. The second step was to interpolate the resulting Rmax for all nodes at each time step of the simulation, to determine the wind velocity throughout the domain at each time step.

The first step was performed in a pre-processing program for all the data from a particular forecast/advisory, before the ADCIRC run began. This design provided visibility to the Rmax values that ADCIRC would use, and gave the Operator the capability to modify the input values for experimentation.

The representative R_{max} values were determined for each quadrant and forecast increment by subtracting the translation speed of the storm from the maximum wind speed, and converting both the translation-reduced maximum wind speed and the highest isotach wind speed in that quadrant into the equivalent wind speeds at the top of the atmospheric boundary layer. These wind speeds and the distance to the highest isotach in that quadrant were then substituted into the gradient wind equation. The gradient wind equation was then solved for the R_{max} in that quadrant using Brent's method, or if that failed numerically, a brute force marching algorithm.

The pre-processing program then appended the resulting R_{max} values to the meteorological input file for use in the actual ADCIRC simulation.

The second step occurred during the execution of ADCIRC. For each time step in the simulation, the central pressure, latitude and longitude of the storm center, and radii to maximum winds were interpolated in time to reflect the simulation time relative to the forecast increments provided by the National Hurricane Center. The translation speed from the most recent forecast increment was used to reduce the time-interpolated maximum wind speed. The time-interpolated maximum wind radii were interpolated in space using a cubic spline; the relevant value of R_{max} at each node was determined from the cubic spline curve. Finally, the Holland(1980) model was used to determine the nodal wind velocity using the spline-interpolated R_{max} , the translation-reduced value of the maximum wind speed, and a Holland B whose value was calculated and then limited to the range of 1.0 to 2.5.

After the computation of the nodal wind velocities at the top of the atmospheric boundary layer, the magnitudes were reduced to corresponding values at 10m, and then reduced again from the 1 minute averaging used by the National Hurricane Center to 10 minute averaging required by ADCIRC. A "damped" translational velocity was vectorially added to the nodal wind velocities, where the damping was proportional to the distance from the radius to maximum winds.

Finally, the wind vectors were rotated toward the center of the storm by an angle that depended on the distance from the center: the rotation angle was ten degrees between the center and the radius to maximum winds; it was twenty five degrees beyond 1.2 times the radius to maximum winds; and the rotation angle was linearly interpolated from the ten degree value at R_{max} and the 25 degree value at 1.2 times R_{max} .

2.5 Output Visualization, Publication, and Interaction

The production of human-comprehensible output is arguably the most important step in the entire process. Clients must be notified that new results are available; and they must be able to get to and use the results in a way that is intuitive for them. As there is more than type of end user or client, there is more than one approach to producing useful output: (a) publication of raw data; (b) production and publication of static images, non-interactive animations, and results files in domain specific formats; and (c) elucidation via interactive visualization through a web application.

Publication of raw results is the most basic step for post processing, and an OpenDAP server was selected and used to publish raw data in NetCDF format to clients and end users. An OpenDAP server provides a web interface to the data, making it easy for users to simply click on a link to the graphics or data file they wish to download, either for the latest run or for any previous simulation run. Technologies such as NCTOOLBOX are available for sophisticated end users to apply their own analyses to the raw data, producing the output they require locally. Once the data have been posted to the OpenDAP server, the system sends an email to a list of email addresses as specified in the system configuration file to notify them that new results are available.

The next level of presentation is in-situ post processing, that is, running non-interactive graphics generation programs in the high performance computing environment to generate static images, non-interactive animations, and reformatted versions of results in specialty formats. These techniques are all in currently place, and have the advantage of not requiring voluminous quantities of data to be moved to an external server for post processing. For example, the system produces GIS shape files, JPEG images and non-interactive animations of maximum inundation as well as Google Earth (kmz) images of maximum inundation, all with in-situ post processing. These output products are then published to the OpenDAP server to make them available to clients. The disadvantage of in-situ processing is the lack of interactivity.

The highest and by far the most effective level of presentation is the use of an interactive web visualization application to elucidate results. For this level, the Coastal Emergency Risks Assessment (CERA) interactive web application was deployed and used to provide an interactive interface that served the needs of non-technical as well as experienced clients.

The CERA application integrates visualization ADCIRC and ADCIRC+SWAN results with Google Maps to provide the context for the results as well as practical features such as panning and the ability to zoom in more closely at various areas of interest to see greater resolution and detail. The application organizes results by storm and advisory (for tropical cyclone results) or by cycle (for NCEP NAM results). It presents these options concisely via dropdown boxes across the top of its interface.

The right side of the interface is an accordion-type menu that presents the various types of data that are available, including water surface elevation, inundation, significant wave height, wave period, and wind speed. For each type of data, the application is able to present a zoomable image of the maximum values that occurred over the course of the forecast (e.g., high water marks) as well as a zoomable, stoppable animation that illustrates the evolution of the data through time.

Furthermore, the interface allowed the user to click on any point of the storm track to view the data that are relevant to the time when the storm was at that point in the track. It is also capable of presenting the nodes of the underlying ADCIRC mesh, and displaying detailed summary information for any particular node the user selects.

The CERA application works by using the OpenDAP server as a staging area for raw data. The CERA application received notification email from the NCFS that new results were published to the OpenDAP server in NetCDF format. It would then download these data as well as some summary information about the type of forecast run that produced the data. It generated visualizations via production of image tiles at progressively finer scales, thus providing the user with the ability to zoom in and examine results in greater detail. These image tiles were then stored on the web server where the data were published, and sent to web clients as required by the user interaction with the web application and its menu system.

3 Configuration

The ASGS is set up for a particular modeling situation using a configuration file. Information in the ASGS configuration file is generally restricted to parameters that are determined by the scientific and logistical choices of the Operator. In contrast, configuration that is particular to the underlying HPC platform is covered in the ASGS Developers Guide, and platform-specific configuration is generally contained within the ASGS itself, rather than an external configuration file.

The ASGS continually re-reads its configuration file; it is first read upon startup, and then once before every nowcast and forecast. This gives the ASGS a degree of dynamic configurability.

One unique aspect of the ASGS configuration file is that it is also a shell script. As a result, the configuration parameters are actually (bash) shell script variables, and are set as such. The use of a shell script as a configuration file also comes in handy if the desired configuration requires some logic, or the use of simple arrays; examples of both will be provided in the following sections.

On the other hand, one of the side effects of the use of a shell script as a configuration file is that it is dangerous to comment out the variables in the configuration file, or to set them equal to "". This causes the associated value to be "missing" in the list of arguments to various subroutines within the ASGS, which moves all the other command line arguments up one place, causing many issues.

In order to avoid these issues, the solution is to leave unneeded configuration parameters alone, rather than commenting them out or setting them to "". If the values are not useful, or not needed, the ASGS will generally deal with them gracefully. There aren't any features or behavior of the ASGS that are triggered by the mere presence of a particular configuration parameter in the configuration file. Parameters can be set to the word "null" for lack of any appropriate value; examples of this are provided below.

Each of the following sections details the parameters that are present in the configuration file. The sections start with a brief description of the section, followed by a snippet from a real configuration file. The definitions of the parameters in that section are then provided in detail. The definitions may be followed by a discussion of how the configuration parameters are used by the ASGS.

In an actual ASGS configuration file, the configuration parameters can be provided in any order.

3.1 Fundamental

Fundamental configuration includes the name of the ASGS instance, resources required, and the type and timing of system startup.

```

INSTANCENAME=2           # "name" of this ASGS process
COLDSTARTDATE=2012042500 # calendar year month day hour YYYYMMDDHH24
HOTORCOLD=coldstart      # "hotstart" or "coldstart"
LASTSUBDIR=null          # path to previous execution (if HOTORCOLD=hotstart)
HINDCASTLENGTH=45.0      # length of initial hindcast, from cold (days)
REINITIALIZESWAN=no      # used to bounce the wave solution

```

DEFINITIONS

INSTANCENAME

The instance name can be any string; a single digit integer is used in practice. The instance name is helpful when multiple instances of the ASGS are running on the same machine. It can also be used to label results, to indicate which instance (among several on a particular machine) generated the results.

COLDSTARTDATE

This date corresponds to ADCIRC time=0; it is used in many different ways by the ASGS, but the principal uses are as follows: (a) determination of nodal factors and equilibrium arguments for tidal forcing; (b) determination of the relationship between the current simulation state (along with the hotstart time) and the available input data.

HOTORCOLD

This parameter can either be set to *coldstart* or *hotstart*. When set to *coldstart*, the ASGS will start by performing a hindcast as described in the previous section. If it is set to *hotstart*, the ASGS will hotstart the simulation using the hotstart file found in the location indicated by the *LASTSUBDIR* parameter described below.

LASTSUBDIR

The full path to the last good nowcast (or hindcast) subdirectory must be provided here if the ASGS is meant to start by resuming a simulation that is already in progress. For example, if the full path to the hotstart file is `/srv/asgs99999/initialize` then the *LASTSUBDIR* parameter should be set to `/srv/asgs99999/initialize`. If the full path to the hotstart file is `/srv/asgs88888/2012010100/nowcast/fort.67`, the *LASTSUBDIR* parameter should be set to `/srv/asgs88888/2012010100`. In other words, the ASGS wants to find either a hindcast or nowcast subdirectory under the *LASTSUBDIR* directory. The ASGS always sets ADCIRC to write a fort.67 file (not a fort.68 file), and only on a hindcast or nowcast, never on a forecast. If the *HOTORCOLD* parameter is set to *coldstart*, this parameter can be set to `null`.

HINDCASTLENGTH

The length of the hindcast is only used by the ASGS if it actually has to perform the hindcast; that is, if the ASGS is supposed to coldstart ADCIRC.

REINITIALIZESWAN

The reinitialization of SWAN parameter can be set to "yes" to re-coldstart SWAN (not ADCIRC) at the start of the nowcast in case the SWAN solution develops issues but the ADCIRC solution has not. Set this back to "no" when the SWAN solution has normalized, or else ASGS will continue to re-coldstart SWAN on every nowcast.

3.2 Source File Paths

The ASGS is flexible about the locations of the files that it requires. This makes it easy to change the version of ADCIRC that is used in the ASGS, for example.

```
ADCIRCDIR=/projects/ncfs/apps/adcirc/trunk/work # ADCIRC executables
INPUTDIR=/projects/ncfs/apps/asgs/trunk/input   # grid and other input files
OUTPUTDIR=/projects/ncfs/apps/asgs/trunk/output # post processing scripts
PERL5LIB=/projects/ncfs/apps/asgs/trunk/PERL    # DateCale.pm perl module
SCRIPTDIR=/projects/ncfs/apps/asgs/trunk        # ASGS executables
```

DEFINITIONS

ADCIRCDIR

The ADCIRC (or ADCIRC+SWAN) directory parameter should be set to the full path where the executable files are located. The use of this parameter allows the Operator to simply compile ADCIRC and leave the executables in place. It also allows the Operator to change the version of ADCIRC that is being used by simply changing the path contained in this parameter.

INPUTDIR

The input directory parameter must be set to the full path to the ADCIRC input files (fort.13, fort.14) and templates (fort.15.template, fort.26.template) that are used by the ASGS to construct input sets, as well as queue script template files that are used to submit compute jobs. When ASGS creates a reusable archive of preprocessed input files, it places it in the directory. This is normally set to the input subdirectory of the ASGS installation.

OUTPUTDIR

The output directory parameter must be set to the full path to the post processing scripts and related data that will be used to post process the ADCIRC results produced by the ASGS. This is normally set to the output subdirectory of the ASGS installation.

PERL5LIB

The perl 5 library parameter is used to set the path to the directory containing the Pcalc.pm perl module. This module is an external dependency, and allows the ASGS to perform the date math that is required for many different purposes during operation. This file is distributed with the ASGS in the PERL subdirectory of the ASGS installation.

SCRIPTDIR

The script directory refers to the ASGS installation itself. There is no assumed relationship between the directory where the ASGS configuration file is stored, or the directory where the ASGS is executed, and the ASGS installation directory. This parameter also allows the Operator to change the version of the ASGS that is used for a particular scenario, by changing this parameter to point to another ASGS installation.

3.3 Physical Forcing

The physical forcing that is applied via ADCIRC is controlled via the following configuration parameters; these parameters are each set to `on` or `off`.

```
BACKGROUNDMET=on      # NAM download/forcing
TIDEFAC=on             # tide factor recalc
TROPICALCYCLONE=off    # tropical cyclone forcing
WAVES=on               # wave forcing
VARFLUX=on             # variable river flux forcing
```

DEFINITIONS**BACKGROUNDMET**

Set to background meteorology to `on` to activate meteorological forcing with data from the North American Meoscale (NAM) model from the National Centers for Environmental Prediction (NCEP). Data are acquired from the NCEP ftp site. If this is turned on, then *TROPICALCYCLONE* must be set to `off` (i.e., these types of forcing are mutually exclusive).

TIDEFAC

Set the tidal factor parameter to `on` to activate tidal forcing, which causes the nodal factors and equilibrium arguments to be recalculated for every simulation run.

TROPICALCYCLONE

The tropical cyclone forcing parameter should be set to `on` to activate meteorological forcing with one of the vortex models that are embedded in ADCIRC. This requires forecast/advisory data from the National Hurricane Center (NHC) as input. Data may be acquired from the NHC web and ftp sites, or may be retrieved from the local file system, as described in the Tropical Cyclone section below. If this is turned on, then *BACKGROUNDMET* should be turned off (i.e., these types of forcing are mutually exclusive).

WAVES

The coupling of SWAN to ADCIRC is activated by setting this parameter to `on`. Because SWAN simulates wind-driven waves, either *BACKGROUNDMET* or *TROPICALCYCLONE* must be turned on for the parameter to have any effect.

VARFLUX

The variable flux parameter is used to activate aperiodic river boundary forcing, and will cause the ASGS to look to an external ftp site for ADCIRC fort.20 files and use that data to construct ADCIRC fort.20 files that are appropriate for the simulation run that it is constructing.

3.4 Computational Resources

The computational resources available to the ASGS will dictate many of the parameter settings in this section, and these parameter settings are interdependent, as described below.

This section also provides an example of the use of shell scripting logic to set some of the parameter values based on other parameter values, to provide a degree of dynamic configurability.

```
TIMESTEPSIZE=0.5          # adcirc time step size (seconds)
SWANDT=1200               # swan time step size (seconds)
NCPU=180                  # number of compute CPUs for all simulations
HINDCASTWALLTIME="24:00:00" # hindcast wall clock time
ADCPREPWALLTIME="00:15:00" # adcprep wall clock time, including partmesh
NOWCASTWALLTIME="05:00:00" # longest nowcast wall clock time
FORECASTWALLTIME="05:00:00" # forecast wall clock time
#QUEUENAME=batch          # default queue name on this machine
#SERQUEUE=batch           # default queue name (for serial jobs)
QUEUENAME=armycore        # dedicated queue name on this machine
SERQUEUE=armycore         # dedicated queue name for serial jobs
#
# the number of processors per node must be specified in the
# queue script for this platform; the number of processors per
# node is different for the different queues; the total number
# of CPUs that we request should be evenly divisible by the
# number of processors per node; lets apply some logic:
if [[ $QUEUENAME = batch ]]; then # this is the default
    PPN=8
    NUMWRITERS=4
else                               # armycore nodes
    PPN=12
    NUMWRITERS=12
fi
```

DEFINITIONS

TIMESTEPSIZE

The time step size (in seconds) is used to set the time step on ADCIRC simulations. A shorter time step generally means that the required wall clock time will be longer, all other things being equal.

SWANDT

The SWAN time step (in seconds) is used to set the time step for SWAN. It should be evenly divisible by the ADCIRC time step size. A shorter timestep generally means a longer wall clock time requirement.

NCPU

The number of CPUs that should be used for all simulations, including hindcast, nowcasts, and each forecast ensemble member.

HINDCASTWALLTIME

The estimated time required to complete the hindcast, whose length (in simulated days) is specified by the *HINDCASTLENGTH* parameter.

ADCPREPWALLTIME

The estimated time required to run adcprep, for the longest expected adcprep job type (i.e., partmesh, prepall, prep15).

NOWCASTWALLTIME

The estimated time required to run a nowcast, taking into account the fact that a nowcast simulation time period may be longer than the typical intercycle simulation time (typically 6 hours) if previous nowcasts have been missed for some reason.

FORECASTWALLTIME

The estimated time required to run a single member of the forecast ensemble, for the ensemble member with the longest forecast length.

QUEUENAME

The name of the queue for parallel computational jobs; this is only present in the configuration file if it varies according to the circumstances of the simulation.

SERQUEUE

The name of the queue for serial (single processor) computational jobs, like adcprep; this is only present in the configuration file if the choice of queue varies according to the circumstances of the simulation.

PPN

The number of processors per node is a parameter that is only required by certain HPC platforms, and even on those platforms it is generally a static parameter set in asgs_main.sh. It is only supplied in the configuration file if it has to be changed depending on the circumstances of the simulation.

NUMWRITERS

If dedicated writer processors are to be used in the simulation, the number of dedicated writers is set here. When submitting a compute job, the ASGS will request a number of processors equal to $N_{CPU} + NUMWRITERS$. If dedicated writers are not to be used, set this value to 0.

3.4.1 Wall Clock Time Estimates

The estimates of wall clock time are required because of the need to fill in the estimated wall clock time of the job in the associated queue script when jobs are submitted. Wall clock times are in the format HH:MM:SS, with a leading zero if the number of hours is a single digit.

On computational platforms where ASGS does not have access to a dedicated queue and is unable to give its jobs special priority over other jobs, the estimated wall clock times should be as small as possible while ensuring that there is enough cushion so that the job is very likely to finish in the estimated time. This is because longer estimated wall clock times generally result in longer waits in a shared queue, with subsequent delays in the production of results.

However, as bad as it is to wait in the queue, it is even worse to have a job killed by the queueing system for exceeding the estimated wall clock limit ... especially if the job was nearly complete at the time it was killed by the system.

In general, an accurate estimate of the wall clock time requires experience with the underlying compute platform as well as the mesh and input files in question. There are many factors that affect wall clock time: number of nodes in the mesh, time step sizes, number of CPUs, implicit vs explicit solution mode in ADCIRC, scalability of the underlying hardware, use of dedicated writer processes, presence or absence of swan coupling, and the type of output files that are being written as well as their frequency of output.

3.4.2 Dynamic Configurability

For this particular HPC platform, there are two queues: a shared queue called "batch" and a high priority queue called "armycore". The platform requires that the number of processors per node be specified, but this parameter depends on the queue to which the job was submitted.

In order to be able to dynamically switch an instance of the ASGS from a low priority queue to a higher priority queue on this platform (for example, if a tropical cyclone is suddenly deemed to be a sufficient threat), the Operator need only change the *QUEUENAME* from *batch* to *armycore*, and script takes care of the other settings. The ASGS then propagates the changes in the queue name, the number of processors per node, and the overall number of CPUs requested into the queue scripts that it submits.

Although the overall number of requested CPUs would change in this scenario, the number of compute CPUs would not. Changing the number of compute CPUs dynamically is also possible, unless ADCIRC is set to produce subdomain hotstart files, or wave coupling has been activated. This is because SWAN always produces subdomain hotstart files, which make it difficult to dynamically change the number of compute CPUs, since the SWAN hotstart files would have to be recomposed into a full domain SWAN hotstart file, and then redecomposed into the right number of subdomain hotstart files for the new number of CPUs.

3.5 External Data Sources

During real time operation, the ASGS contacts external sites over the network to acquire the data it needs for its next nowcast/-forecast cycle. The configurability of this contact allows the Operator to create mock sites for testing purposes, as well as keep up with changes in the structure of external data sites.

3.5.1 Tropical Cyclones

```
STORM=03                # storm number, e.g. 05=ernesto in 2006
YEAR=2010               # year of the storm
TRIGGER=rss             # either "ftp" or "rss"
RSSSITE=www.nhc.noaa.gov # site information for retrieving advisories
FTPSITE=ftp.nhc.noaa.gov # hindcast/nowcast ATCF formatted files
FDIR=/atcf/afst         # forecast dir on nhc ftp site
HDIR=/atcf/btk          # hindcast dir on nhc ftp site
#
#RSSSITE=www.seahorsecoastal.com # used for testing
#FTPSITE=ftp.unc.edu          # for test storms
#FDIR=/pub/ims/rick/ADCIRC/NHC_Advisories/fst # forecast dir on test site
#HDIR=/pub/ims/rick/ADCIRC/NHC_Advisories/btk # hindcast dir on test site
```

DEFINITIONS

STORM

The storm parameter is the two digit storm number assigned by the National Hurricane Center. The name of storm can change, e.g., TD TWELVE can become HU KATRINA, but the number never changes (still 12 in that example). The NHC names data files according to the storm number, not its name.

YEAR

The four digit year is also used by the NHC to name data files that are relevant to a particular tropical cyclone.

TRIGGER

The trigger parameter tells the ASGS where to look for new advisories. When set to `rss` or `rsseembedded`, the ASGS looks at the advisory number in an RSS feed (i.e., an `index-at.xml` file); when the number changes, a new advisory is deemed to have been issued. In the past, the ASGS always followed a link in the RSS feed to a website with the text of the new forecast/advisory. However, starting with the 2012 season, the NHC has also started embedding the text of the forecast/advisory directly in the xml of their RSS feed. In addition, they sometimes don't update the link to the new forecast/advisory page when they update the advisory number; this issue caused 2 stale advisories to be fed to the ASGS at the start of the 2012 season (1 during Alberto and 1 during Beryl). As a result, the ASGS can now be configured to pull the text of the forecast/advisory directly from the RSS feed, rather than following the embedded link. This feature is activated by setting `TRIGGER=rsseembedded`. In order to configure the ASGS to follow the RSS feed's link to the text of the forecast/advisory, Operators should set `TRIGGER=rss`. In the past, it was also possible to set the `TRIGGER` to `ftp`, causing the ASGS to look for a changed forecast file on the NHC ftp site. As a result of a variety of issues associated with getting forecast/advisory data in ATCF format from the NHC ftp site, `TRIGGER` should not be set to `ftp` for tropical cyclone events.

RSSSITE

The Really Simple Syndication (RSS) site is the website where the ASGS should look for the `index-at.xml` (the "at" indicates that it is the feed for the Atlantic basin). If `RSSSITE` is set to "filesystem", then `FDIR` must be set to the full path to the `index-at.xml` (RSS feed) file on the local filesystem.

FTPSITE

If `FTPSITE` is set to "filesystem", then `HDIR` must be set to the path to the `bal<storm><year>.dat` (ATCF-formatted) hindcast/nowcast data file.

FDIR

The forecast directory is the path to the `index-at.xml` on the website where advisories are being issued if `RSSSITE` is a website. If `RSSSITE` is set to `filesystem`, then the forecast directory should be set to the full path to the `index-at.xml` file on the local filesystem.

HDIR

The hindcast directory is the path to the bal<STORM><YEAR>.dat file containing the hindcast/nowcast data for the storm *STORM* of year *YEAR* on the ftp site *FTPSITE*. If *FTPSITE* is set to `filesystem` then the hindcast directory should be set to the full path to the bal<STORM><YEAR>.dat file on the local filesystem.

Testing

New for the 2012 season, the ASGS can now be configured to load tc hindcast files and/or forecast/advisory files directly from the local filesystem, instead of downloading them from remote web and ftp sites.

This feature is useful for testing, and for production on platforms that do not support direct download of external files. It is activated by setting *FTPSITE* to `filesystem` and/or setting *RSSSITE* to `filesystem` in the ASGS configuration file. The *FTPSITE* and *RSSSITE* configuration parameters are completely independent of each other.

Interactions

If the *RSSSITE* is set to `filesystem` and *TRIGGER* `rss`, then the ASGS will load the index-at.xml from the designated location on the local file system and parse out the link to the text of the forecast advisory. If it is a link to a website, it will follow the link to the website.

However, if the forecast/advisory is also on the local filesystem, the <link> tag in the RSS xml should be populated with the full path and file name of the corresponding forecast/advisory text file.

When *TRIGGER* is set to `rssemdedded`, the ASGS still looks at the advisory number in the RSS feed to determine when a new advisory has been issued, but instead of following a hyperlink, it pulls the text of the forecast/advisory from the RSS feed itself.

When *TRIGGER* is set to `ftp`, the ASGS will look for an ATCF-formatted file called al<stormnumber><year>.fst, and compare the content of that file to the last one it downloaded. When the content changes, a new advisory is deemed to have been issued.

The problems with setting *TRIGGER* to `ftp` include the following: (a) the ATCF-formatted file does not contain the advisory number; (b) the NHC sometimes issues intermediate advisories simply to update nowcast information, which will look to the ASGS like a new advisory; (c) the file can contain typos which do not affect the ASGS but the typo will be corrected by NHC, causing the ASGS to think a new advisory has been issued.

On the other hand, the ASGS still uses `ftp` to get the ATCF-formatted BEST track files (that is, the hindcast/nowcast information); it now also has the option of looking to the local filesystem instead of a remote ftp site.

Since *TRIGGER*=`ftp` is not used, *FDIR* is only used by the ASGS when *RSSSITE*=`filesystem`, to find the RSS feed (index-at.xml) on the local file system.

3.5.2 Background Meteorology

```
FORECASTCYCLE="00,12"
BACKSITE=ftp.ncep.noaa.gov          # NAM forecast data from NCEP
BACKDIR=/pub/data/nccf/com/nam/prod # contains the nam.yyyymmdd files
FORECASTLENGTH=84                  # hours of NAM forecast to run (max 84)
PTFILE=ptFile_oneEighth.txt        # the lat/lons for the OWI background met
ALTNAMDIR="/projects/ncfs/data/asgs5463", "/projects/ncfs/data/asgs14174"
```

DEFINITIONS**FORECASTCYCLE**

The forecast cycle parameter tells the ASGS which NAM cycles to run a forecast for, and which NAM cycles it should just nowcast. It will always download forecast files for every cycle in any case. Set this to "00,06,12,18" to run the forecast for all NAM cycles. Set this to `none` to only run nowcasts (i.e., if forecasting should be turned off).

BACKSITE

The background meteorological ftp site is the location where the ASGS should look for background meteorological data.

BACKDIR

The background meteorological directory is the location on *BACKSITE* where the directories containing NAM output files in grib2 format are located (directory naming convention is nam.yyyymmdd).

FORECASTLENGTH

Number of hours to forecast; default value is the total length of the NAM forecast (84 hours or 3.5 days).

PTFILE

The point file contains a list of lat/lon points that the NAM data should be reprojected to during the conversion from grib2 format and Lambert Conformal projection to OWI (ascii text) format and geographic projection. The utility *input/pt-File_gen.pl* is provided with the ASGS to aid in generating or regenerating files of this type.

ALTNAMDIR

These local directories are alternate locations for searching for NAM input data on the local filesystem, in case old data cannot be found on the NCEP ftp site. When looking for nowcast files, the ASGS will look for grib2 files in the directories \$ALTNAMDIR/*/*nowcast/erl.YYMMDD/" (2 digit year, 2 digit month, 2 digit day). When looking for forecast data, the ASGS will look for grib2 files in the following directories: \$ALTNAMDIR/YYYYMMDDHH24/namforecast/ (4 digit year, 2 digit month, 2 digit day, 2 digit hour of 24 hour clock). The directories must be separated by commas with no spaces. No need for a trailing forward slash.

3.5.3 River Flux

If *VARFLUX* has been set to on, then the ASGS will download aperiodic river flux data (i.e., adcirc fort.20 files from an external website, as configured in this section.

```
RIVERSITE=ftp.nssl.noaa.gov
RIVERDIR=/projects/ciflow/adcirc_info
```

DEFINITIONS**RIVERSITE**

The river boundary data ftp site parameter provides the domain name of the anonymous ftp site that the ASGS should contact for fort.20 files with variable river flux data.

RIVERDIR

The river directory parameter is the full path to the river flux files on the *RIVERSITE* anonymous ftp site.

3.6 Input Files and Templates

The input files and input file templates are used by the ASGS to construct simulation runs. Many of these files are not included in the ASGS repository, as many of them are very large. If the files to be used are not in the repository, they will have to be provided by the Operator.

These parameters are used to specify just the names of the input files and templates; the full path is not provided. The ASGS will look for these files on the filesystem path specified in the *INPUTDIR* parameter.

```
GRIDFILE=nc_inundation_v6c_rivers_msl.grd # mesh (fort.14) file
GRIDNAME=nc6b
CONTROLTEMPLATE=v6brivers_nowindreduction_explicit_fort.15_template # fort.15 template
ELEVSTATIONS=v6brivers_elev_stations.txt
VELSTATIONS=null
METSTATIONS=v6brivers_met_stations.txt
NAFILE=v6brivers_newrough.13
SWANTEMPLATE=fort.26.v6b.limiter.template
RIVERINIT=v6brivers.88
RIVERFLUX=v6brivers_fort.20_default
HINDCASTRIVERFLUX=v6brivers_fort.20_hc_default
PREPPEDARCHIVE=prepped_ncv6b_${INSTANCENAME}_${NCPU}.tar.gz
HINDCASTARCHIVE=prepped_ncv6b_hc_${INSTANCENAME}_${NCPU}.tar.gz
```

DEFINITIONS**GRIDFILE**

The grid file parameter specifies the name of the mesh file (fort.14 file) to be used in the simulation.

GRIDNAME

The grid name is just a short and human-readable name to identify the mesh; it is used in postprocessing to label the output.

CONTROLTEMPLATE

The control template is an adcirc fort.15 file (control file) that has been converted into a template by replacing key input values with special characters. During execution, the ASGS looks for the special characters and replaces them with the appropriate values for an actual simulation run. More information about this is provided in the ASGS Developers Guide.

ELEVSTATIONS

The elevation stations file is a list of the elevation recording stations that should be output to the adcirc fort.61 file. The file format is the same as it would be for a list of stations in the adcirc fort.15 file.

VELSTATIONS

The velocity stations file is like the elevation stations file, except it is for current velocity and the adcirc fort.62 file.

METSTATIONS

The meteorological stations file is like the elevation stations file, except it is for meteorological data and the adcirc fort.71 and fort.72 files.

NAFILE

The nodal attributes file parameter provides the name of the nodal attributes file (fort.13 file) that should be used with the simulation. If there is no nodal attributes file, this parameter can be left blank.

SWANTEMPLATE

The swan template file is an adcirc fort.26 file (swan control file) that has been converted into a template using the same type of process as for the *CONTROLTEMPLATE* described above.

RIVERINIT

If *VARFLUX* has been set to on, an upland river initialization file (adcirc fort.88 file) must be provided to provide the initial water level heights along the river nodes.

HINDCASTRIVERFLUX

If *VARFLUX* has been set to on, an adcirc fort.20 file (aperiodic river flux boundary file) must be provided to cover the entire duration of the hindcast (as specified by the *HINDCASTLENGTH* parameter). The hindcast river flux parameter must be set to the name of this fort.20 file.

PREPPEDARCHIVE

Because of the delay caused by having to run adcprep (specifically partmesh and prepall) on large meshes, the ASGS will create an archive of the subdomain fort.13, fort.14, and fort.18 files to avoid having to run these time consuming adcprep jobs more than once. The use of this prepped archive allows the ASGS to only use adcprep for fort.15 and fort.20 files on each simulation run, since those are the only files that change from cycle to cycle. The number of compute CPUs is embedded in the file name in the example above, as this allows dynamic configurability of the number of compute CPUs (if the number of compute CPUs changes, the ASGS will not find the prepped archive that it needs, and will generate a new one for the new number of CPUs. Please note that **if any of the station files are changed, the prepped archive file will no longer be valid and must be deleted**. If the prepped archive file does not exist, the ASGS will detect this and simply generate a new one.

HINDCASTARCHIVE

The prepped hindcast archive file is created for the same purpose as the *PREPPEDARCHIVE* file but is only used on a hindcast. The reason for a separate archive for hindcasts is that meteorological forcing is not applied on a hindcast, so meteorological stations are not present in the adcirc fort.15 file on a hindcast, which changes the adcirc fort.18 (message passing) files. As a result, the subdomain files for hindcasts must be kept separate from the subdomain files for nowcasts and forecasts, at least when meteorological forcing is applied to nowcasts and forecasts, which is always, for all practical purposes.

Although the ASGS does not generally include the actual mesh and other input files for production instances of the ASGS, the east coast 95d mesh and input file are included for use in learning and testing the ASGS. In order to use these example files, use the following configuration settings:

```

GRIDFILE=ec_95d.grd                                # mesh (fort.14) file
GRIDNAME=ec95d
CONTROLTEMPLATE=ec_95_fort.15_template # fort.15 template
ELEVSTATIONS=corps_elev_stations.txt    # or substitute your own stations file
VELSTATIONS=corps_vel_stations.txt
METSTATIONS=corps_vel_stations.txt
NAFILE=                                           # no nodal attributes for ec95d
SWANTEMPLATE=fort.26.ec95.template        # only used if WAVES=on
RIVERINIT=null                                 # this mesh has no rivers ...
RIVERFLUX=null
HINDCASTRIVERFLUX=null
PREPPEDARCHIVE=prepped_${GRIDNAME}_${INSTANCENAME}_${NCPU}.tar.gz
HINDCASTARCHIVE=prepped_${GRIDNAME}_hc_${INSTANCENAME}_${NCPU}.tar.gz

```

3.7 Storm Ensemble

Once a nowcast run is complete, the ASGS is capable of running more than one type of forecast. If background meteorology is used, the forecast ensemble can only consist of one member: `namforecast`. However, if tropical cyclone forcing is used, the ensemble can have one or more members, with each member expressing a different perturbation from the base storm.

```

ENSEMBLESIZE=1 # number of storms in the ensemble
STORMLIST[0]=8 # namforecast
NAME[0]="nowcast"
NAME[1]="nhcConsensus"
NAME[2]="higherMaxWindSpeed"
NAME[3]="slowerOverlandSpeed"
NAME[4]="veerRight"
NAME[5]="veerLeft"
NAME[6]="largerRmax"
NAME[7]="smallerRmax"
NAME[8]="namforecast"
PERCENT[2]=20
PERCENT[4]=100
PERCENT[5]=-100
PERCENT[6]=50
PERCENT[7]=-50
#ENDTIME=2008090200

```

DEFINITIONS

ENSEMBLESIZE

The ensemble size parameter should be set to the number of different forecast runs to be performed after each nowcast. For background meteorology (NAM), or for tropical cyclone forcing where only the NHC consensus forecast should be run, the ensemble size should be set to 1. For tropical cyclone forecasting where multiple perturbations to the base forecast are to be run, the ensemble size should be equal to the total number of forecast runs.

STORMLIST

The storm list is an array parameter of size *ENSEMBLESIZE* that provides the array index of the name of each storm in the ensemble. The ASGS is sensitive to these names, and uses them during construction of input files to apply the specified perturbations.

NAME

The name list is an array parameter of arbitrary size that lists the names of the various storms that could potentially be used in an ensemble. In other words, this array will generally have lots of values that may or may not be used in any particular storm ensemble.

PERCENT

The percent variation is an array parameter of size *ENSEMBLESIZE* that specifies the percent variation for perturbations that are implied by the storm name. The ASGS ignores this parameter for storms named *nowcast*, *nhcConsensus*, and *namforecast*.

ENDTIME

The optional end time parameter provides the year, month, day, and hour when the forecast period should end. This parameter was implemented for storms whose forecast reaches so far inland that it becomes irrelevant; cutting the forecast short saves processing time for landfalling storms and allows the ASGS to produce results more quickly.

The *STORMLIST*, *NAME*, and *PERCENT* parameters represent examples of the ASGS configuration file taking advantage of its identity as a shell script by using of simple bash shell array variables as configuration parameters. All of the arrays are zero indexed.

3.8 Output Files

Control over ADCIRC output is mainly accomplished through the use of a single parameter, *OUTPUTOPTIONS*. The content of this parameter is one long string of command line options (ultimately fed to the *control_file_gen.pl* perl script) that specify the characteristics of the various output files.

The *OUTPUTOPTIONS* parameter is built by concatenating several convenience parameters for each of the output file types. These convenience parameters include *FORT63*, *FORT7374*, *SPARSE*, etc; these sub-parameters are really only used to enhance readability of the configuration file. It would be equally valid (but less readable) to supply the same command line options directly in the *OUTPUTOPTIONS* parameter.

By default, all output from ADCIRC is turned off; output of each type can be activated by providing the corresponding command line argument for the output frequency. In order for output to be produced, the output frequency must be nonzero.

If a particular type of output is activated, then a new output file of that type will be created for each nowcast/forecast cycle (rather than appending the file from the previous nowcast/forecast cycle) by default. On the other hand, if continuous output is desired, this behavior can be specified with the corresponding "append" string (e.g., *--fort61append*), which will cause the data file to reflect the whole time series from the time that the append parameter was supplied.

Output frequency in this section is in SECONDS, not timesteps; the ASGS takes care of calculating the ADCIRC timestep increment required to achieve the output frequency specified here, and will automatically recalculate the timestep increment if the Operator changes the ADCIRC time step using the *TIMESTEPSIZE* parameter.

By default, output files are in ADCIRC ascii text format. The format of output files can be specified on a file-by-file basis, by providing the corresponding netcdf format string (e.g., *--fort63netcdf*).

```
FORT61="--fort61freq 900.0 --fort61netcdf" # water surface elevation station output
FORT62="--fort62freq 0" # water current velocity station output
FORT63="--fort63freq 3600.0" # full domain water surface elevation output
FORT64="--fort64freq 3600.0" # full domain water current velocity output
FORT7172="--fort7172freq 3600.0" # met station output
FORT7374="--fort7374freq 3600.0" # full domain meteorological output
SPARSE="--sparse-output"
OUTPUTOPTIONS="${SPARSE} ${FORT61} ${FORT62} ${FORT63} ${FORT64} ${FORT7172} ${FORT7374}"
HOTSTARTCOMP=fulldomain # fulldomain or subdomain hotstart files
HOTSTARTFORMAT=netcdf # binary or netcdf hotstart files
MINMAX=reset # "continuous" or "reset" for maxele.63 etc fil
```

DEFINITIONS

OUTPUTOPTIONS

The output options parameter is a quoted string containing the command line options to the `control_file_gen.pl` perl script for controlling output. Its content is the amalgamation of other subparameters, if any.

FORT61

The fort.61 sub-parameter is optional, and is used to control output to the ADCIRC water surface elevation recording stations (fort.61) file. The locations and names of the recording stations are provided via a file specified by the *ELEVSTATIONS* parameter (see the Input Files and Templates section of this document).

FORT62

The fort.62 sub-parameter is optional, and is used to control output to the ADCIRC water current velocity recording stations (fort.62) file. The locations and names of the recording stations are provided via a file specified by the *VELSTATIONS* parameter (see the Input Files and Templates section of this document).

FORT63

The fort.63 sub-parameter is optional, and is used to control output to the ADCIRC fulldomain water surface elevation (fort.63) file.

FORT64

The fort.64 sub-parameter is optional, and is used to control output to the ADCIRC fulldomain water current velocity (fort.64) file.

FORT7172

The fort.71/fort.72 sub-parameter is optional, and is used to control output to the ADCIRC meteorological recording stations files (fort.71 and fort.72) file. The locations and names of the recording stations are provided via a file specified by the *METSTATIONS* parameter (see the Input Files and Templates section of this document).

FORT7374

The fort.73/fort.74 sub-parameter is optional, and is used to control output to the ADCIRC fulldomain meteorological output (fort.73 and fort.74) files.

SPARSE

The sparse output parameter is optional, and if provided, it will cause all fulldomain non-netcdf output files for which ADCIRC is capable of producing sparse output to be produced in that format.

HOTSTARTCOMP

The composition of hotstart files can be either "subdomain" to produce subdomain hotstart files, or "fulldomain" to produce fulldomain hotstart files. This parameter only affects ADCIRC hotstart files; SWAN always produces subdomain hotstart files. This parameter has no effect if the ADCIRC hotstart files are in netcdf format, as ADCIRC is not capable of producing subdomain hotstart files in netcdf format.

HOTSTARTFORMAT

The hot start format parameter is used to specify whether hotstart files are produced in nonportable binary or in netcdf format. Because of the transparency and portability of NetCDF hotstart files, the NetCDF format is recommended.

MINMAX

The maxele.63 and maxwvel.63 files are always generated by ADCIRC and will normally reflect the solution since cold-start, if ADCIRC was coldstarted, or the solution since the most recent hotstart if ADCIRC was hotstarted. If this parameter is set to `reset`, these two files will not be copied from the previous nowcast cycle, so that these files will only reflect the current nowcast/forecast cycle, "forgetting" the high water marks etc from previous cycles. If this parameter is set to `continuous`, then the ASGS will copy these files from the previous nowcast/forecast cycle to the current cycle, preserving the history of extreme values. **Care should be taken in the setting of the MINMAX value to `continuous`, including consultation with end users, who may be confused by the inclusion of high water marks from the past when these high water marks are presented in the context of a forecast, which end users expect to only contain results pertaining to the future.**

3.9 Notification

The ASGS is capable of sending email to interested parties at several points in its execution, to notify them of its status or of the occurrence of various events (both positive and adverse).

Parameters that are populated with more than one email address should have each address separated by a space; the whole list of addresses should be enclosed in double quotes.

The email address lists should be kept as short as possible; if a particular email address is receiving email at each stage of the ASGS execution, this could result in more than a dozen emails per day, evenly spaced throughout the day and night.

Experience shows that jittery end users do not want to hear their smart phone rattling on their night stand at 4:30am with a message from the ASGS that the NHC has just issued a new advisory, but the results will not be available for another couple hours.

```
EMAILNOTIFY=yes          # yes to have host HPC platform email notifications
NOTIFY_SCRIPT=ncfs_nam_notify.sh
ACTIVATE_LIST="foo@bar.edu foo2@bar.com"
NEW_ADVISORY_LIST="foo@bar.edu foo2@bar.com"
POST_INIT_LIST="foo@bar.edu foo2@bar.com"
POST_LIST="foo@bar.edu foo2@bar.com"
JOB_FAILED_LIST="foo@bar.edu foo2@bar.com"
NOTIFYUSER=foo@bar.edu
ASGSADMIN=foo@bar.edu
```

DEFINITIONS

EMAILNOTIFY

The email notify parameter can be set to `yes` to enable the ASGS to send out notifications via email; if set to `no`, all email notification is turned off, except for email sent to the ASGS Operator directly from the queueing system when a job fails.

NOTIFY_SCRIPT

The notify script parameter is used to specify the shell script with the content of the actual emails that will be sent in each situation. Because different ASGS sites have different requirements for the content of this email, its content is configurable by simply plugging in a different script and selecting it via this parameter. The notification script is expected to be found in the output subdirectory of the *SCRIPTDIR* directory.

ACTIVATE_LIST

The activate list parameter contains one or more email addresses of recipients that should receive an announcement when the ASGS is

NEW_ADVISORY_LIST

The new advisory list parameter contains one or more email addresses of recipients that should receive notification that a new advisory is now running.

POST_INIT_LIST

The post processing initialization list contains one or more email addresses that should receive notification that a new forecast is complete and that post processing has started.

POST_LIST

The in-situ post processing completion list contains one or more email addresses to receive notification that the in-situ post processing (that is, post post processing and graphics generation that are set to occur locally on the HPC platform where ADCIRC is running) is complete, and the results have been made available for viewing, interaction, and/or downloading (depending on the site and type of output product).

JOB_FAILED_LIST

The job failed list parameter contains one or more email addresses that receive notification when the ASGS detects that a job has failed.

NOTIFYUSER

The notify user parameter contains one email address, that of the ASGS Operator. This email address is propagated by the ASGS to the queue scripts for the computational jobs that it submits. As a result, the Operator will receive email directly from the queueing system on the HPC platform where the ASGS is running. Queueing systems generally only send email when a job experiences an error. This parameter is not affected by the *EMAILNOTIFY* parameter above; the ASGS Operator will always receive these types of emails, even if *EMAILNOTIFY* is set to `no`.

ASGSADMIN

The ASGS administrator parameter contains one email address, that of the ASGS Operator. This address receives an email if the ASGS experiences a fatal error; the content of the email is the ASGS log file up to the time that the fatal error occurred.

3.10 Post Processing and Publication

The ASGS post processing infrastructure includes in-situ post processing (i.e., generation of graphics and/or other output products on the HPC platform where the ASGS itself is running) as well as publication of results to end users. It is one area of ASGS implementation that can be wildly different at different sites. Some ASGS sites produce all their graphics and other output products in-situ, while other sites simply copy raw ADCIRC output files to an external server.

This diversity is accommodated by asking the Operator to specify the name of an executable program to run at the end of a forecast that performs the desired post processing tasks. Changing the post processing and publication then becomes as easy as changing the name of the post processing program in this configuration file.

```
INITPOST=null_init_post.sh
POSTPROCESS=ncfs_post.sh
POSTPROCESS2=null_post.sh
TARGET=blueridge
WEBHOST=webserver.hostingco.com
WEBUSER=remoteuser
WEBPATH=/home/remoteuser/public_html/ASGS/outputproducts
```

DEFINITIONS

INITPOST

The initial post processing parameter specifies the name of an executable program that the ASGS should run at the start of a forecast. This may be used to create directories for the results on an external website. To turn off initial post processing, set this parameter to the trivial script `null_init_post.sh`, which does nothing.

POSTPROCESS

The post process parameter specifies the name of an executable program that the ASGS should use for in-situ post processing and publication of results. See the existing post processing scripts like `ncfs_post.sh` and `corps_post.sh` for examples. See the code in `asgs_main.sh` to see the list of arguments that are supplied to this post processing script. To turn off post processing, set this parameter to the value `null_post.sh` which is a trivial shell script that does nothing.

POSTPROCESS2

If a second, separate post processing script is required, it can be incorporated using this parameter; the associated executable is run by the ASGS after running the executable indicated by the `POSTPROCESS` parameter.

TARGET

The target parameter is specific to the programs in RenciGETools package that generate Google Earth (kmz) visualizations; it is used to select the target geographical location for the visualizations from a predefined list. See the RenciGETools subdirectory (included with the ASGS) for details.

WEBHOST

The optional web host parameter is useful for specifying the hostname of a machine with a webserver that will publish results; output products will be transferred via scp to this host, provided that ssh key authentication has been configured.

WEBUSER

The optional web user parameter specifies the user name on the host `WEBHOST` that should be used in establishing the scp connection between the HPC platform where ASGS is running and the remote host.

WEBPATH

The optional web path parameter is the full path to the directory where output products should be transferred on `WEBHOST`.

The ASGS assumes that the executables indicated by *INITPOST*, *POSTPROCESS*, and *POSTPROCESS2* are located in the *OUTPUTDIR* directory.

One important issue to take into account in the design of the post processing is error handling. Specifically, the post processing script developer should ensure that control is returned to the calling routine (*asgs_main.sh* in this case) in the event of an error. Under no circumstances should the post processing script enter a loop that waits on a particular subprocess to complete successfully; if that subprocess ends abnormally, an infinite loop will result, preventing control from returning to the ASGS and preventing any further nowcast/forecast cycles from running.

3.11 Archiving

Once results have been produced, there is an issue of how they should be stored. This is another area, like post processing, that varies widely among sites that run the ASGS. Some HPC platforms have specific policies regarding data retention, and well developed systems in place for long term storage. Other sites have an ad-hoc policy regarding data retention and no long term storage facilities of any kind. Again, the ASGS handles this diversity by allowing the Operator to plug in an executable file that performs the desired tasks.

```
ARCHIVE=ncfs_archive.sh
ARCHIVEBASE=/projects/ncfs/data
ARCHIVEDIR=archive
```

DEFINITIONS

ARCHIVE

The archive parameter is the name of an executable file that is called by the ASGS after all forecast ensemble members are complete. It is automatically executed in the background (that is, with an ampersand) so that the ASGS can continue with the next nowcast/forecast cycle immediately, if needed. The ASGS passes the values of *ARCHIVEBASE* and *ARCHIVEDIR* to the archival executable, along with other input arguments (see *asgs_main.sh* for details). The ASGS expects to find the executable program specified by this parameter in the *OUTPUTDIR* directory.

ARCHIVEBASE

The archive base parameter represents the full path to the location where files should be archived.

ARCHIVEDIR

The archive directory parameter represents the subdirectory of *ARCHIVEBASE* where files should be archived. It will be created by the ASGS if it does not already exist.

4 Execution

How to make it go, using screen, starting in the right directory, and how to check the log files.

5 Testing

5.1 Setting Up the Mock NHC Files for Testing

For the forecast files, you can use the text advisories from the NHC website.

<http://www.nhc.noaa.gov/archive/>

Just click on the year, and you'll see all the storms for that year.

Then click on the name of the storm you want to run. You'll see all the forecast/advisories that were issued for that storm.

Then click on the forecast/advisory you want to run, and "Save Page As..." to save the html (just the html) to a file. You could feed this html file to the ASGS, but we are going to do some additional editing, as described below. Repeat this step for all the advisories you want to run.

Once you have those, you'll need a mock RSS feed. You can create this by taking the text of each forecast/advisory from the html files you downloaded and inserting it into an index-at.xml file. I've attached the index-at.xml file from Beryl advisory 07 as an example.

For example, if you want to run Ike advisories 42, 44, and 46, you'd do the following

1. use a text editor to pluck out the actual text of the forecast/advisory from each html file;
2. use a text editor to embed that text into a template index-at.xml file, being sure to also update the advisory number;
3. end up with three files named something like ike.42.index-at.xml, ike.44.index-at.xml, ike.46.index-at.xml.

Now, you'll also need a hindcast file to go along with each forecast file. In a real storm, the hindcast data end before or just after the start of the forecast when each advisory is issued.

So, in the case of Ike, you need a set of bal092008.dat files that end at the appropriate times for your forecasts. To support our test efforts in prior seasons, I've created a bunch of files like this for various storms here:

```
ftp://ftp.unc.edu/pub/ims/rick/ADCIRC/NHC_Advisories/btk/
```

Hopefully this archive will be of use to you.

So at this point, you have a bunch of files called (for example) ike.XX.index-at.xml and ike_advisory_XX.btk (one set for each advisory, where XX is the advisory number).

5.2 Configuring the ASGS for Test Advisories Stored on Local Filesystem

```
STORM=  
YEAR=  
TRIGGER=rsseembedded  
RSSSITE=filesystem  
FTPSITE=filesystem  
FDIR=/path/to/stormname.XX.index-at.xml  
HDIR=/path/to/stormname_advisory_XX.btk
```

Using Ike as an example, would be 09 and would be 2008.

5.3 Actually Running ASGS from Advisories Stored in Local Files

1. Go to the directory where you have the mock RSS feed for the forecast/advisories.
2. Make a symbolic link to the first advisory you want to run, e.g.:

```
ln -s ike.42.index-at.xml index-at.xml
```

because the ASGS is hardwired to look for the index-at.xml file.

1. Make a symbolic link to the hindcast data that correspond to the first advisory you want to run, e.g.:

```
ln -s ike_advisory_42.btk bal092008.dat
```

because the ASGS is hardwired to look for a file that is formatted as bal.dat.

1. Start up the ASGS as you normally would. It will grab the NHC data from the local filesystem (whichever advisory your symbolic links point to).
2. Whenever you want to issue a new advisory, manually update the symbolic link for the hindcast to the next hindcast file, then manually update the symbolic link to the next forecast file. For example:

```
rm bal092008.dat ; ln -s ike_advisory_44.btk bal092008.dat
rm index-at.xml ; ln -s ike.44.index-at.xml index-at.xml
```

Always update the symbolic link to the hindcast first, b/c the ASGS will be looking for a change in the forecast data to tell it when a new advisory has been issued. If you update the forecast link first, the ASGS will detect it so quickly that it will start its new cycle b/f you can update the hindcast link. This will just cause it to use the old hindcast.

You could even write a shell script to update these links periodically if you want; it doesn't have to be manual, although it probably will be when you first start doing this.

6 Performance During Hurricane Irene

Hurricane Irene 2011 provided an opportunity to use the ASGS and CERA technologies to assist officials in the NWS and USCG with real-life operational and emergency management decisions. The implementation of these technologies for the North Carolina and US East Coast that provided the basis for this real life application are described below, including the static physical data, use of dynamic meteorological and river flow data, operational configuration, interaction with web-based results, and feedback from official end users.

During Hurricane Irene of 2011, the ASGS and CERA technologies were actively and heavily utilized by National Weather Service (NWS) offices, emergency managers, and the US Coast Guard. In one example of the real benefit for emergency management officials, US Coast Guard Admirals personally operated the web application to examine and interact with simulation results, and then used those results in their decision to evacuate their command center near Norfolk, VA, relocating to their secondary command center in Missouri. Subsequently, the predicted flooding did occur, rendering the abandoned USCG command center useless. In summary, the USCG was able to maintain uninterrupted command and control over the course of Hurricane Irene from their alternate command center, and have lauded the ASGS/CERA system for its contribution to this successful operational decision.

6.1 North Carolina Physical Data

The discretization of the North Carolina domain (the mesh) that was used for both ADCIRC and SWAN is shown below in figure 3. The mesh contained 295328 nodes and 575512 with a minimum element size of 13m. The land cover data that were used to derive directional wind roughness lengths, canopy coefficients, and Manning's n friction values are shown in figure 5. The system set 127 points throughout the domain for sampling the model output for comparison with tide gauges and water levels in river flood plains and other locations. It also specified 99 locations for point sampling of model meteorological data for comparison with physical meteorological data collected at land based and buoy-based meteorological data collected by various state and federal agencies and universities.

6.2 Dynamic Input Data

As the storm that would eventually become Irene was forming, the ASGS/CERA system had already been running for months at RENCi on the blueridge machine, a 1408 core Dell High Performance Computing cluster, producing daily forecasts using NCEP NAM meteorology, river flow data from NSSL, and tidal forcing.

When Irene formed and the decision was made to start running vortex-forced tropical cyclone forecasts, a new instance of the ASGS was started up from the most recent hotstart file from the existing NAM results. The new instance was configured to run on 384 processors, rather than the 180 that were used in the NAM runs. Both ASGS instances then ran side by side for several days, sending their output to the CERA web application, until the NAM instance was shut down to conserve processor time.

6.3 Internally Generated Input Files

On 384 processors, the ASGS was able to perform all file manipulations for the nowcast and forecast cycle for the consensus track, submit the jobs, and generate results for a 5 day forecast in 1 hour 15 minutes from the time the NHC forecast/advisory was issued.

As hurricane Irene approached the North Carolina coast, the team member that operates of the ASGS system lost electric power and could therefore no longer oversee the operation of the ASGS during this critical period. However, since the ASGS itself was actually executing at the RENCi facility, which is located in Durham, North Carolina and that facility never lost power, the ASGS was able to continue operating autonomously without interruption and without interaction or oversight for the duration of the storm.

6.4 Output Visualization and Interaction

The results produced by the ASGS were converted in-situ to Google Earth (kmz) graphics and posted to the OpenDAP server at RENCi for use by the Newport NWS forecast office. The raw data in NetCDF format were also posted to the OpenDAP server and then a notification email was sent to the CERA web application.

The web application picked up the latest raw data, generated the image tiles for the animations and maximum values visualizations at all levels of detail, and sent email notification to selected official end users at NOAA and the USCG to notify them that new results were available for interaction over the web. The CERA website itself was not password protected, but end users were instructed not to publicize the availability of the site. The CERA web processing required another 1 hour and 15 minutes to generate image tiles for a particular advisory, for a total turn around time of 2 hours and 30 minutes.

There were myriad examples of the use of the ASGS/CERA results by emergency managers and officials during the storm event. As stated in the executive summary, several Admirals at the USCG personally examined and interacted with the ASGS/CERA results via the CERA site and subsequently made the operational decision to evacuate their command center in Virginia. The flooding depicted in the results did occur, and the USCG was extremely pleased and satisfied with their experience with the ASGS/CERA guidance.

7 Performance During the Deepwater Horizon Event

Discuss configuration, operation, and performance during Deepwater Horizon.

8 Refs

“Irene helps UNC researchers hone forecasts” UNC Spotlight, <http://www.unc.edu/spotlight/Irene-tests-researchers>, retrieved 18 March 2012.

“Hurricane Drives Toward New York with Deadly Fury”, New York Times Sunday Edition, August 28, 2011; scientific credit for real time results used in a graphic depicting storm surge guidance along US east coast.

Fleming, J., C. Fulcher, R. Luettich, B. Estrade, G. Allen, and H. Winer, 2008. A Real Time Storm Surge Forecasting System using ADCIRC, Estuarine and Coastal Modeling X, M. L. Spaulding [ed], ASCE, 373-392.

A This Document

This document was prepared from the text file ASGSOperatorsGuide.txt using software called asciidoc (<http://www.methods.co.nz/asciidoc/>). The document can be formatted as an html page with the command

```
asciidoc ASGSOperatorsGuide.txt
```

or formatted as a pdf with the command

```
a2x --format=pdf ASGSOperatorsGuide.txt
```